

MAML-Select: An Online Adaptive Client Selection Method for Federated Learning via Meta-Learning

Abstract—In Federated Learning (FL), choosing which clients update the shared model in each round is a combinatorial scheduling problem under resource constraints. Clients differ in data utility, training speed, battery capacity and bandwidth. Since only a small cohort can communicate per round, conventional client-selection strategies often favor clients with high training utility or fast updates, which can repeatedly select a narrow group of clients, leaving slower or temporarily low-utility clients underused, reducing participation coverage, and making training process less efficient. Existing methods like FedCS and TiFL use system state to reduce straggler effects, whereas Oort, FedCor, CriticalFL, and FedGCS score clients using training utility, loss correlation, critical-period timing, or learned client combinations. However, these methods often rely on a fixed scoring rule, or update it slowly, even though a client’s usefulness changes as the global model evolves. Model-Agnostic Meta-Learning Select (MAML-Select) removes this fixed-rule assumption, and frames each round as a fast adaptation task. A lightweight policy network is learned using the previous round’s observed loss reduction and latency, and is used for ranking clients. After the selected clients return their updates, the same policy is meta-updated using the new feedback. We evaluate MAML-Select on Fashion-MNIST, CIFAR-10, and CIFAR-100 with 100 clients, $K = 10$ selected clients per round, a Dirichlet non-independent and identically distributed (non-IID) split with $\alpha = 0.5$, and three random seeds. On Fashion-MNIST, CIFAR-10, and CIFAR-100, MAML-Select provides competitive accuracy with significantly lower cumulative Tera Floating-Point Operations (TFLOPs) and modelled energy consumption. MAML-Select also retains full participation coverage on every dataset.

Impact Statement—FL is central to privacy-preserving AI across heterogeneous environments, such as mobile health, intelligent transportation, industrial sensing, and smart-city systems, but these deployments often face tight computational-capacity, energy-consumption, and latency budgets. MAML-Select contributes an adaptive selection mechanism that reduces unnecessary client-training computational cost while preserving broad participation coverage. By using recent training feedback to re-rank clients, the method supports FL systems that respond to changing device conditions instead of relying on a fixed rule. The reported energy-consumption values are modelled estimates that follow the declared simulator and device-tier setup, giving a transparent basis for comparing selection policies. The broader impact is a more resource-aware path for deploying distributed AI at the edge, including the industrial Internet of Things (IIoT), especially when model updates must remain efficient, reproducible, and privacy-preserving.

Index Terms—Federated learning, client selection, meta-learning, and computational efficiency.

I. INTRODUCTION

FEDERATED Learning (FL) trains a shared model across many devices without moving their raw data to the server [1]. Clients train models locally and share only model updates. However, they differ in hardware, network conditions, and local data. Their heterogeneity creates statistical drift under non-independent and identically distributed (non-IID)

data and straggler delays under synchronous training [2]–[6]. Despite these advances, the problem of *whom to select* each round remains open.

Client selection chooses a subset $\mathcal{S}_t \subset \{1, \dots, N\}$ each round. Client utility is non-stationary, because a useful client can become redundant later, while an overlooked client can become important in the long run. System-aware methods improve client participation, but their selection bias can leave slow yet data-rich clients underused [7], [8]. Utility-aware methods prioritize information gain, but may incur high latency or expensive server-side modelling [9]–[13]. Learning-based selectors use histories, features, or reinforcement signals [14]–[16]. An efficient round-by-round mechanism for adapting the ranking policy to evolving utility remains under-explored.

We propose MAML-Select, a client-selection framework based on a Model-Agnostic Meta-Learning (MAML)-style fast adaptation step [17], which learns the policy network via recent feedback before the next ranking decision. Unlike reinforcement-learning (RL)-based or contextual-bandit selectors, the proposed MAML-Select uses rapid gradient-based policy adaptation. Unlike personalization of FL client models [18], the proposed MAML adapts the selector, rather than the client model. The proposed approach is useful when devices have limited computational capacity, strict deadlines, or intermittent connectivity. The major contributions can be summarized as follows.

- 1) We formulate the FL client selection problem as a meta-learning task and introduce the MAML-Select algorithm with rapid gradient-based adaptation for online client selection.
- 2) Experiments on Fashion-MNIST, CIFAR-10, and CIFAR-100 under non-IID settings ($\alpha = 0.5$, three seeds) show that MAML-Select delivers accuracy competitive with FedAvg and FedGCS at lower computational cost, and higher accuracy than FedCS, Oort, TiFL, and FedCor. Sensitivity and ablation studies over λ , the inner-loop steps, the selector width, and the state features support the design choices.
- 3) We demonstrate that MAML-Select *lowers cumulative training TFLOPs* by up to 21.5% and modelled energy consumption by up to 24.7% relative to FedAvg, with paired-test support for the resource reductions.

Organization. The rest of this paper is organized as follows. Section II reviews related work on client selection in FL. Section III details the system model and problem formulation. The proposed MAML-Select algorithm is described in Section IV, followed by evaluation and results in Section V. Section VI concludes this paper and discusses future research directions.

II. RELATED WORK

In the following, the classification of existing client selection methods for the FL algorithm are presented.

1) *System-Aware Selection*: System-aware selection addresses stragglers. FedCS selects feasible clients under resource constraints, while TiFL groups devices into online performance tiers [7], [8]. Recent studies show similar concerns for slow but data rich clients in resource-constrained medical IoT and asynchronous FL with heterogeneous objectives [4], [5]. MAML-Select uses feedback from the previous training phase, and learns the selection policy based on six parameters.

2) *Data-Aware and Utility-Based Selection*: Data-aware methods select clients for model utility rather than only speed. Oort prioritizes data-rich and fast clients, and FedCor models correlations between client losses [9], [10]. CriticalFL uses critical learning periods, while FedGCS searches for client combinations in a learned continuous space [12], [13]. Other work couples selection with aggregation design [11], [19]. MAML-Select instead targets rapid adaptation of the ranking policy as client utility evolves.

3) *Learning-Based and Meta-Learning Approaches*: Learning-based selectors improve from observed feedback. FAVOR uses a deep Q-network, while neural contextual bandits learn from client features and rewards [14], [15], and recent work studies deep-RL selection for robustness and fairness-aware participation [16], [20]. MAML is used in [18] to personalize client models, whereas MAML-Select adapts the server-side ranking policy instead.

Table I consolidates these distinctions. For methods without a stated closed-form selection complexity, it reports the dominant server-side operation.

Difference from prior work. MAML-Select meta-learns the selection policy h_ϕ . Unlike heuristic, RL, contextual-bandit, and generative selectors, it uses a MAML inner update to convert recent feedback into fast policy weights before the next cohort is scored. A fixed heuristic or bandit may update rewards over time, but does not adapt the scoring-network parameters from a support set immediately before selection.

III. SYSTEM MODEL AND PROBLEM FORMULATION

We consider a central server and N heterogeneous clients indexed by $\mathcal{N} = \{1, 2, \dots, N\}$. Client i stores a local dataset $\mathcal{D}_i$ with n_i samples. The global objective is to minimize the empirical risk over the distributed data as

$$\min_{\theta \in \mathbb{R}^d} F(\theta) = \sum_{i=1}^N \frac{n_i}{n} F_i(\theta), \quad (1)$$

where $\theta \in \mathbb{R}^d$ denotes the global model, $n = \sum_{i=1}^N n_i$, and $F_i(\theta) = \frac{1}{n_i} \sum_{j=1}^{n_i} \ell(\theta; x_{i,j}, y_{i,j})$ is the local empirical risk, with $\ell(\cdot)$ the per-sample loss (e.g. cross-entropy).

At round t , let $\mathcal{A}_t \subseteq \mathcal{N}$ be the available clients and $\mathcal{S}_t \subseteq \mathcal{A}_t$ the selected cohort. The server holds the global model and the policy, while clients keep data local and return only updates, loss summaries, and resource-state feedback, so selection is a server-side decision taken before each broadcast.

Client i has per-round latency $T_{i,t} = T_{i,t}^{comp} + T_{i,t}^{comm}$. The computation time is $T_{i,t}^{comp} = E n_i C_{model} / f_{i,t}$, where

$E n_i C_{model}$ is the local-training work over E epochs and n_i samples, C_{model} is the per-sample forward-plus-backward cost in FLOPs, and $f_{i,t}$ is the device processing rate set by its hardware tier. The communication time is $T_{i,t}^{comm} = M / B_{i,t} + d_{i,t}$, for an update of size M , uplink bandwidth $B_{i,t}$, and fixed delay $d_{i,t}$ [7]. In synchronous FL the slowest selected client sets the round latency,

$$T_t^{round} = \max_{i \in \mathcal{S}_t} T_{i,t}. \quad (2)$$

Let $a_{i,t}$ indicate whether client i is selected. Under policy π_ϕ , the constrained client-selection problem is

$$\begin{aligned} \min_{\pi_\phi} \quad & \sum_{t=1}^T [F(\theta_{t+1}) + \lambda T_t^{round}] \\ \text{s.t.} \quad & \sum_{i=1}^N a_{i,t} = K, \quad a_{i,t} \in \{0, 1\}, \\ & a_{i,t} \leq \mathbf{1}\{i \in \mathcal{A}_t\}, \end{aligned} \quad (3)$$

where K is the cohort size, T is the number of rounds, and $\lambda \geq 0$ controls the trade-off between latency and loss, and the client selection subset is updated as $\mathcal{S}_t = \{i : a_{i,t} = 1\}$. A larger λ places more emphasis on latency, while $\lambda = 0$ yields loss-only selection. We assume $|\mathcal{A}_t| \geq K$. The model θ_{t+1} results from local training and aggregation over $\mathcal{S}_t$. The problem remains challenging because the action space is combinatorial and client utility evolves during training.

IV. MAML-SELECT ALGORITHM

Building on (3), MAML-Select treats client selection as a meta-learning problem, because client utility changes as training progresses. The algorithmic details of the selection policy are given below.

1) *State Space*: The server builds a state vector $\mathbf{s}_{i,t} \in \mathbb{R}^6$ for each client i at round t . It holds the local training loss $\hat{L}_{i,t}$, the gradient norm $\|\nabla F_i\|_2$, the estimated latency $\hat{T}_{i,t}$, the battery level $b_{i,t}$, the selection frequency $\nu_{i,t}$, and the staleness $\tau_{i,t}$ measured in rounds since the client was last selected [11], [15]. The three pairs capture statistical utility, system feasibility, and participation balance in turn. Each round t is a task $\mathcal{T}_t$ in which clients are ranked so that the top- K cohort minimizes the per-round cost.

The six raw entries carry different physical units. Before they reach the policy, each column is therefore standardized across the clients available in that round,

$$\mathbf{s}_{i,t} \leftarrow (\mathbf{s}_{i,t}^{raw} - \boldsymbol{\mu}_t) \oslash (\boldsymbol{\sigma}_t + \epsilon_0), \quad (4)$$

where $\boldsymbol{\mu}_t, \boldsymbol{\sigma}_t$ are the per-feature mean and standard deviation over $\mathcal{A}_t$, $\oslash$ is elementwise division, and $\epsilon_0 = 10^{-8}$. The policy input is then dimensionless and the score is a relative ranking within one round, which is all Top- K needs.

MAML-Select trains a meta-policy network $h_\phi(\mathbf{s}_{i,t})$, parameterized by ϕ , which is adapted to the current round using recent feedback. The policy is a 6-64-64-1 multilayer perceptron (MLP) with rectified linear unit (ReLU) activations and 4,673 trainable parameters. Cost feedback for a cohort becomes observable only after that cohort has finished training, so the two meta-learning sets are drawn from two consecutive

TABLE I
COMPARISON OF REPRESENTATIVE METHODS AND MAML-SELECT.

Method	Primary objective	Adaptation mechanism	Selection complexity or dominant operation
FedCS [7]	Increase feasible updates under resource constraints	Resource estimates before selection	Resource-constrained subset search
TiFL [8]	Mitigate stragglers while preserving accuracy	Online tier updates	Tiering and probabilistic sampling
Oort [9]	Improve time-to-accuracy	Utility and speed feedback	Utility scoring and guided sampling
FedCor [10]	Accelerate convergence under non-IID data	Loss-correlation model	Gaussian process (GP) covariance operations
FAVOR [14]	Counterbalance non-IID bias	Experience-driven policy updates	DQN inference and training
NCCB [15]	Learn contextual client combinations	Client features and reward feedback	Locality-sensitive hashing (LSH) feature extraction and bandit updates
FedCG [11]	Couple selection and gradient compression	Capability-aware compression updates	Submodular optimization and linear programming
CriticalFL [12]	Exploit critical learning periods	Period-aware participation	Base-selector dependent
FedGCS [13]	Optimize client combinations	Search in a learned continuous space	Latent-space optimization and beam search
Kazemi et al. [16]	Improve robustness under poisoning attacks	Experience-driven policy updates	Deep-RL inference and training
Vox Populi [20]	Promote fair and inclusive participation	Fairness-aware client selection	Framework-dependent
Per-FedAvg [18]	Personalize the client model to local data	MAML adaptation of the client model, not the selector	No client selection; per-client model fine-tuning
MAML-Select	Adapt rankings under evolving client utility	Fast adaptation from recent round feedback	$\mathcal{O}(N \phi)$ scoring plus a gradient-adaptation step

completed rounds. Writing $\mathcal{S}_t$ for the cohort of round t , the support and query sets used at round t are

$$\begin{aligned}\mathcal{D}_{sup}(t) &= \{(\mathbf{s}_{i,t-2}, c_{i,t-2})\}_{i \in \mathcal{S}_{t-2}}, \\ \mathcal{D}_{query}(t) &= \{(\mathbf{s}_{i,t-1}, c_{i,t-1})\}_{i \in \mathcal{S}_{t-1}},\end{aligned}\quad (5)$$

and therefore $\mathcal{D}_{sup}(t) = \mathcal{D}_{query}(t-1)$. This lag is what makes the procedure a meta-learning problem rather than online regression, because the inner step adapts on an older round while the meta-gradient is measured on a newer one. The outer update therefore rewards initializations from which one adaptation step transfers forward in time, which is what a selector needs when utility drifts. The candidate set $\mathcal{D}_{cand}(t) = \{\mathbf{s}_{i,t} \mid i \in \mathcal{A}_t\}$ carries no labels and is only scored.

2) *From the Cohort Objective to a Per-Client Score:* The objective in (3) is defined over cohorts, whereas the selector ranks clients individually. We connect the two in two steps, and only the second is inexact.

First, the straggler term is a maximum and is therefore not separable. For any deadline $T_{target} > 0$ and any cohort $\mathcal{S}_t$,

$$\frac{T_t^{round}}{T_{target}} \leq 1 + \sum_{i \in \mathcal{S}_t} \rho_{i,t}, \quad \rho_{i,t} = \frac{(T_{i,t} - T_{target})_+}{T_{target}}, \quad (6)$$

where $(x)_+ = \max(0, x)$. The left side is at most 1 when every selected client meets the deadline, and otherwise the slowest client j contributes $1 + \rho_{j,t} \leq 1 + \sum_i \rho_{i,t}$. The bound never understates the latency cost, makes it additive, and renders the penalty dimensionless so it can be traded against a loss reduction through one unitless weight λ .

Second, the risk decrease $F(\theta_t) - F(\theta_{t+1})$ is unknown when the cohort is chosen, so we replace it by $\sum_{i \in \mathcal{S}_t} \Delta_{i,t}$, the observable local loss reductions $\Delta_{i,t} = \hat{L}_{i,t}^{before} - \hat{L}_{i,t}^{after}$ over the E local epochs. This is the one approximation in the derivation. It treats cohort contributions as additive and equates local with global progress, ignoring client drift and

interference between updates, and we adopt it because it needs no validation set and no extra communication.

Substituting (6) and this surrogate into the per-round term of (3) and dropping the additive constant λ gives a cohort surrogate that is modular in $\mathcal{S}_t$,

$$\tilde{F}_t(\mathcal{S}) = \sum_{i \in \mathcal{S}} c_{i,t}, \quad c_{i,t} = \lambda \rho_{i,t} - \Delta_{i,t}, \quad (7)$$

which is the per-client cost used throughout this work and the quantity the policy is trained to predict. Both terms are dimensionless, since $\rho_{i,t}$ is a fractional deadline overrun and $\Delta_{i,t}$ a loss difference. The deadline T_{target} is the mean expected round latency of the Tier-2 devices in the pool, recomputed every round, and the pool median is used if no Tier-2 device is available.

Proposition 1 (Top- K is exact for the surrogate). *Let $\mathcal{S}^*$ hold K indices of $\mathcal{A}_t$ with the smallest $c_{i,t}$. Then $\mathcal{S}^*$ minimizes $\tilde{F}_t$ over all subsets of $\mathcal{A}_t$ of size K .*

Proof. Suppose some feasible $\mathcal{S}$ has a smaller objective. Since both sets have size K , there are $u \in \mathcal{S} \setminus \mathcal{S}^*$ and $v \in \mathcal{S}^* \setminus \mathcal{S}$ with $c_{v,t} \leq c_{u,t}$. Swapping u for v does not increase the objective and strictly reduces $|\mathcal{S} \setminus \mathcal{S}^*|$, so at most K swaps turn $\mathcal{S}$ into $\mathcal{S}^*$ without ever increasing it, a contradiction.

The chain is therefore an exact bound, one stated approximation, and an exact combinatorial step. One gap remains, since $c_{i,t}$ depends on quantities realized only after client i has trained, and the meta-policy fills it by predicting $c_{i,t}$, after which Proposition 1 applies to the predicted costs. We claim only that the cohort exactly optimizes a stated modular surrogate under predicted costs.

A. The MAML-Select Algorithm

The algorithm proceeds in three phases per round, all executed at the start of round t in the order below. Both meta-learning updates therefore complete before any client is scored, and they use only the lagged feedback of (5), so no quantity from round t selects the cohort of round t .

Phase 1, inner loop adaptation. At the start of round t , the server adapts the meta-policy ϕ_t on $\mathcal{D}_{sup}(t)$ using one gradient-descent step,

$$\phi'_t = \phi_t - \beta \nabla_{\phi_t} \mathcal{L}_{sup}(h_{\phi_t}; \mathcal{D}_{sup}(t)), \quad (8)$$

where $\mathcal{L}_{sup}$ is the mean-squared-error (MSE) loss between predicted scores and observed costs from round $t-2$, and β is the inner learning rate. The step is plain gradient descent with $\beta = 0.01$ and a single inner step.

Lemma 1 (Inner-step descent). *Let the support loss $g_t(\phi) = \mathcal{L}_{sup}(h_{\phi}; \mathcal{D}_{sup}(t))$ be L -smooth, that is, ∇g_t is L -Lipschitz. If the inner step size obeys $0 < \beta < 2/L$, the inner update (8) is a descent step on the support loss,*

$$g_t(\phi'_t) \leq g_t(\phi_t) - \beta \left(1 - \frac{L\beta}{2}\right) \|\nabla g_t(\phi_t)\|_2^2.$$

Proof. By L -smoothness, g_t obeys the descent lemma

$$g_t(\psi) \leq g_t(\phi) + \nabla g_t(\phi)^\top (\psi - \phi) + \frac{L}{2} \|\psi - \phi\|_2^2.$$

Substituting $\psi = \phi'_t$, $\phi = \phi_t$, and the inner update $\phi'_t - \phi_t = -\beta \nabla g_t(\phi_t)$ yields the stated bound. Its coefficient $\beta(1 - L\beta/2)$ is positive precisely for $0 < \beta < 2/L$, so the step never increases g_t and strictly decreases it unless $\nabla g_t(\phi_t) = 0$.

A single support-set step therefore reduces the round's ranking loss, which justifies adapting before the candidates are scored.

Phase 2, outer loop meta-update. The meta-policy is updated on the query set of (5),

$$\phi_{t+1} \leftarrow \phi_t - \eta \nabla_{\phi'_t} \mathcal{L}_{query}(h_{\phi'_t}; \mathcal{D}_{query}(t)), \quad (9)$$

where η is the meta-learning rate. Query-loss gradient is evaluated at the adapted weights ϕ'_t and applied to the meta-parameters ϕ_t , which omits the second-order derivatives. Because $\mathcal{D}_{query}(t)$ is the feedback of round $t-1$ and $\mathcal{D}_{sup}(t)$ that of round $t-2$, this update measures how well one adaptation step on older feedback predicts the round that followed it. In the implementation the direction $\nabla_{\phi'_t} \mathcal{L}_{query}$ is applied through Adam at $\eta = 10^{-3}$ rather than by plain descent, and Remark 1 states what this means for the analysis.

Phase 3, client selection and FL execution. The adapted policy $h_{\phi'_t}$ scores the available clients and selects the K clients with the lowest predicted costs,

$$\mathcal{S}_t = \text{Top-}K_{i \in \mathcal{A}_t} (-h_{\phi'_t}(s_{i,t})). \quad (10)$$

Two rules qualify (10). During the first $\lceil N/K \rceil$ rounds the policy has too little feedback to rank anything, so a cold start takes the cohort in a fixed seed-shuffled coverage order that visits every client once. Thereafter an exploration slot gives one of the K places to the stalest and least-used available client, ordered by $(-\tau_{i,t}, \nu_{i,t})$, and the rest to the Top- $(K-1)$ of the learned ranking. The slot is deterministic rather than random, is set to $\varepsilon = 1$, and bounds how long any client can stay unobserved. Sec. V-C verifies that coverage does not depend on the cold start.

The selected clients then train, and their states and realized costs become the query set of round $t+1$ and the support set of round $t+2$. Algorithm 1 gives the full procedure.

Lemma 2 (Outer-step descent). *Let the query loss $q_t(\phi) = \mathcal{L}_{query}(h_{\phi}; \mathcal{D}_{query}(t))$ be L_q -smooth, and let*

$$\delta_t = \|\nabla q_t(\phi'_t) - \nabla q_t(\phi_t)\|_2$$

denote the look-ahead displacement, the gap between the direction actually applied, $\nabla q_t(\phi'_t)$, evaluated at the adapted weights ϕ'_t , and the query gradient $\nabla q_t(\phi_t)$ at the un-adapted weights. If the outer step size obeys $0 < \eta \leq 1/L_q$, the meta-update (9) satisfies

$$q_t(\phi_{t+1}) \leq q_t(\phi_t) - \frac{\eta}{2} \|\nabla q_t(\phi_t)\|_2^2 + \frac{\eta}{2} \delta_t^2,$$

with the inner step controlling this error through $\delta_t \leq L_q \beta \|\nabla g_t(\phi_t)\|_2$.

Proof. Write $\phi_{t+1} = \phi_t - \eta d_t$ with $d_t := \nabla q_t(\phi'_t)$. With $\eta \leq 1/L_q$, the descent lemma for q_t gives

$$q_t(\phi_{t+1}) \leq q_t(\phi_t) - \eta \nabla q_t(\phi_t)^\top d_t + \frac{\eta}{2} \|d_t\|_2^2.$$

Completing the square, with $-a^\top b + \frac{1}{2} \|b\|_2^2 = \frac{1}{2} \|b - a\|_2^2 - \frac{1}{2} \|a\|_2^2$ for $a = \nabla q_t(\phi_t)$ and $b = d_t$, rewrites the right-hand side as

$$q_t(\phi_t) - \frac{\eta}{2} \|\nabla q_t(\phi_t)\|_2^2 + \frac{\eta}{2} \|d_t - \nabla q_t(\phi_t)\|_2^2.$$

Since $\|d_t - \nabla q_t(\phi_t)\|_2 = \delta_t$, the first bound follows. The second uses L_q -smoothness and (8), which give $\delta_t \leq L_q \|\phi'_t - \phi_t\|_2 = L_q \beta \|\nabla g_t(\phi_t)\|_2$.

Since the bound makes δ_t proportional to β for a bounded support-gradient norm, the error term $\frac{\eta}{2} \delta_t^2 = \mathcal{O}(\beta^2)$ is small at $\beta = 10^{-2}$, and the first-order meta-update behaves like gradient descent on the query loss up to that $\mathcal{O}(\beta^2)$ floor.

Two clarifications are needed. First, δ_t is not the term dropped by the first-order approximation, which is $\beta \nabla^2 g_t(\phi_t) \nabla q_t(\phi'_t)$. It measures how far the applied direction sits from the query gradient at ϕ_t , which is what the descent inequality needs, and Lemma 2 controls only that. Second, a rate follows only if the query objective is fixed across rounds, which is untenable here. Since $\mathcal{D}_{sup}(t) = \mathcal{D}_{query}(t-1)$, setting $q_t \equiv q$ would force $g_t \equiv q$ and collapse the problem into descent on a fixed loss, and it would assert that client utility never changes, which Sec. V-D shows is false. We therefore state a tracking bound.

Corollary 1 (Tracking under a drifting query objective). *Let Lemma 2 hold at every round with $0 < \eta \leq 1/L_q$, let the query objectives be bounded below by $q_{\min}$, and let $V_T = \sum_{t=1}^T (q_{t+1}(\phi_{t+1}) - q_t(\phi_{t+1}))$ be the cumulative drift, measured at a common iterate. Then*

$$\min_{1 \leq t \leq T} \|\nabla q_t(\phi_t)\|_2^2 \leq \frac{2(q_1(\phi_1) - q_{\min})}{\eta T} + \frac{2V_T}{\eta T} + \frac{1}{T} \sum_{t=1}^T \delta_t^2.$$

Proof. Lemma 2 gives $\frac{\eta}{2} \|\nabla q_t(\phi_t)\|_2^2 \leq q_t(\phi_t) - q_t(\phi_{t+1}) + \frac{\eta}{2} \delta_t^2$. Consecutive terms belong to different objectives, so write $q_t(\phi_t) - q_t(\phi_{t+1}) = [q_t(\phi_t) - q_{t+1}(\phi_{t+1})] + [q_{t+1}(\phi_{t+1}) - q_t(\phi_{t+1})]$. Summing, the first bracket telescopes to at most $q_1(\phi_1) - q_{\min}$ and the second to V_T . Dividing by $\eta T/2$ and bounding the average below by its minimum gives the claim.

Algorithm 1 MAML-Select Algorithm

Require: clients $\mathcal{N}$ with $|\mathcal{N}| = N$, rounds T , cohort size K , inner LR β , outer LR η , trade-off λ , exploration slot ε

- 1: **Initialize** global model θ_0 , meta-policy ϕ_1 and its Adam state, coverage order $\pi \leftarrow$ seeded shuffle of $\mathcal{N}$, and $\mathcal{D}_{sup}, \mathcal{D}_{query} \leftarrow \emptyset$.
- 2: **for** round $t = 1, \dots, T$ **do**
- 3: Broadcast θ_t ; collect states $\mathbf{s}_{i,t}$ from $\mathcal{A}_t$, standardized by (4)
- 4: **if** $\mathcal{D}_{query} \neq \emptyset$ **then**
- 5: $\phi'_t \leftarrow$ inner step (8) on $\mathcal{D}_{sup}$ ▷ Phase 1
- 6: $\phi_{t+1} \leftarrow$ Adam step on ϕ_t along $\nabla_{\phi'_t} \mathcal{L}_{query}(h_{\phi'_t}; \mathcal{D}_{query})$ ▷ Phase 2
- 7: **else**
- 8: $\phi'_t \leftarrow \phi_t, \phi_{t+1} \leftarrow \phi_t$ ▷ no feedback yet
- 9: **end if**
- 10: score $_i \leftarrow h_{\phi'_t}(\mathbf{s}_{i,t})$ for all $i \in \mathcal{A}_t$ ▷ Phase 3
- 11: **if** $t \leq \lceil N/K \rceil$ **then** ▷ cold start
- 12: $\mathcal{S}_t \leftarrow$ next K clients of π , cyclically.
- 13: **else**
- 14: $\mathcal{E}_t \leftarrow \varepsilon$ clients of $\mathcal{A}_t$ largest in $(-\tau_{i,t}, \nu_{i,t})$. ▷ exploration
- 15: $\mathcal{S}_t \leftarrow \mathcal{E}_t \cup \text{Top-}(K - \varepsilon)$ of $-\text{score}$ over $\mathcal{A}_t \setminus \mathcal{E}_t$.
- 16: **end if**
- 17: Receive $\{\Delta\theta_{i,t}\}_{i \in \mathcal{S}_t}$, local losses, and system metrics
- 18: $\theta_{t+1} \leftarrow \theta_t + \sum_{i \in \mathcal{S}_t} \frac{n_i}{\sum_{j \in \mathcal{S}_t} n_j} \Delta\theta_{i,t}$
- 19: $T_{target} \leftarrow$ mean Tier-2 latency over $\mathcal{A}_t$; costs $c_{i,t}$ by (7), $i \in \mathcal{S}_t$
- 20: $\mathcal{D}_{sup} \leftarrow \mathcal{D}_{query}, \mathcal{D}_{query} \leftarrow \{(\mathbf{s}_{i,t}, c_{i,t}) \mid i \in \mathcal{S}_t\}$.
- 21: **end for**
- 22: **return** θ_{T+1}

Remark 1 (Scope of Corollary 1). If the environment settles, so that $V_T = o(T)$, the policy reaches an $\mathcal{O}(\beta^2)$ -approximate stationary point at rate $\mathcal{O}(1/(\eta T))$. A fixed objective gives $V_T = 0$, and any $V_T \leq 0$ recovers the same rate. If utility drifts at a constant rate the drift term does not vanish and the guarantee degrades to tracking, meaning the policy stays within a drift-proportional neighbourhood of stationarity. We therefore make no claim of convergence to a stationary point of a fixed objective. The analysis also assumes an outer step $\phi_{t+1} = \phi_t - \eta \nabla_{\phi_t} q_t(\phi'_t)$, whereas the implementation passes that direction through Adam, so it characterizes the applied direction but not Adam’s per-coordinate rescaling.

Fig. 1(b) is consistent with this reading, since the meta-update magnitude on CIFAR-100 shrinks and then settles at a small non-zero level rather than decaying to zero.

MAML-Select differs from fixed-rule selectors because its policy is adapted before each decision, and from RL selectors because the update is an explicit gradient step on recent feedback.

1) *Computational Complexity:* Let $P = |\phi|$ be the number of meta-policy parameters, so $P = 4,673$ here. Scoring N clients costs $\mathcal{O}(NP)$, a size- K heap selects the best K in $\mathcal{O}(N \log K)$, and the inner and outer updates each touch at most the cohort and cost $\mathcal{O}(KP)$. The per-round server-side selection cost is therefore

$$\mathcal{O}(NP + N \log K + 2KP). \quad (11)$$

To validate (11) we measured the selection overhead over 10 rounds for $N \in \{20, 40, 80, 100, 200, 500, 1000\}$ with the selection rate held at 10%. Fig. 2(b) shows it staying between about 20 and 26 ms per round as N grows fifty-fold, while the operation count in Fig. 2(c) grows about linearly. The measured time stays flat because at these pool sizes the small $\mathcal{O}(NP)$ scoring cost is dominated by fixed per-round bookkeeping. Since P and K are fixed in the protocol, MAML-Select scales linearly with N . The memory cost is $\mathcal{O}(P + N + K)$ for the policy parameters, client scores, and selected cohort. This contrasts with correlation-matrix or

Gaussian-Process selection approaches, where the covariance operations grow beyond linearly in N , and cubically in common FedCor-style implementations [10].

V. EVALUATION AND RESULTS

We evaluate on three benchmarks. Fashion-MNIST uses LightCNN-9 [21], and CIFAR-10 and CIFAR-100 use ResNet18 [22]. We simulate $N = 100$ clients under a Dirichlet non-IID partition at $\alpha = 0.5$, selecting $K = 10$ clients per round for $T = 200$ rounds, or 150 rounds on CIFAR-100, with $E = 5$ local epochs and batch size 32. Table II reports mean $\pm$ standard deviation over seeds 42, 123, and 2026. Local models use PyTorch default initialization and stochastic gradient descent (SGD) with learning rate 0.01, momentum 0.9, weight decay 10^{-4} , and a constant scheduler. The meta-policy uses one inner step with $\beta = 0.01$ and Adam for the outer update with $\eta = 0.001$. The selector seed is fixed to 2026, so the reported variance reflects training and partition randomness alone, and Algorithm 1 with these parameters fully specifies the simulation. Processing rates follow three device tiers reflecting the systems heterogeneity common in federated deployments [23], namely Tier 1 at 20% and $1.0\times$, Tier 2 at 50% and $2.0\times$, and Tier 3 at 30% and $4.0\times$. The same tiers drive resource accounting, with 4 W and 18 Wh batteries for Tier 1, 7 W and 30 Wh for Tier 2, and 12 W and 48 Wh for Tier 3. The deadline T_{target} is the average round latency of Tier 2 clients and $\lambda = 0.5$. Baselines are FedAvg, FedCS, Oort, TiFL, FedCor, FedGCS, and CriticalFL, as characterized in Table I.

A. Accuracy and Resource Cost

We estimate the client-training computational cost per round as

$$\text{TFLOPs}_{total} = \sum_{t=1}^T \sum_{i \in \mathcal{S}_t} E n_i C_{model} \times 10^{-12}, \quad (12)$$

where C_{model} matches $T_{i,t}^{comp}$, so cost and latency stay consistent, and one ResNet18 round on CIFAR-10 is about 4.13 TFLOPs. Modelled energy accumulates per-tier power over local-training time,

$$\mathcal{W}_{total} = \sum_{t=1}^T \sum_{i \in \mathcal{S}_t} \frac{\mathcal{P}_i T_{i,t}^{comp}}{3600}, \quad (13)$$

where $\mathcal{P}_i$ is the device-tier power in watts and $T_{i,t}^{comp}$ is in seconds, so $\mathcal{W}_{total}$ accumulates in watt-hours (Wh).

Table II summarises all methods. MAML-Select matches FedAvg and FedGCS on Fashion-MNIST and CIFAR-100, and exceeds FedCS, Oort, TiFL, FedCor, and CriticalFL on CIFAR-100. On CIFAR-10 it gives up 3.8 pp against FedAvg for substantially lower cost. It therefore offers a resource-aware operating point close to the strongest accuracy baselines, and its meta-update overhead of about 10^{-6} TFLOPs per round is negligible next to the client-training cost in Table II.

Paired t -tests over the three seeds support the resource reductions against FedAvg on every dataset, with TFLOPs p -values of 0.009, 0.002 and 0.012 and energy p -values of 0.007, 0.001 and 0.003. Against FedGCS the accuracies are




Fig. 1. CIFAR-100 convergence over 150 rounds. The panels report (a) accuracy, (b) the selector meta-update magnitude $\|\phi_{t+1} - \phi_t\|_2$ for MAML-Select, which shrinks and then settles at a small non-zero level, the behaviour Corollary 1 predicts under a drifting query objective, (c) cumulative modelled energy consumption, and (d) test loss.

indistinguishable on CIFAR-10 ($p = 0.10$) and CIFAR-100 ($p = 0.49$). Against FedAvg they are equivalent on Fashion-MNIST (-0.1 pp, $p = 0.84$), while CIFAR-10 (-3.8 pp) and CIFAR-100 (-1.5 pp) trade a small accuracy loss for the best energy efficiency. Lowering λ from 0.5 to 0.1 raises CIFAR-10 diagnostic accuracy by only about one point, so the residual gap reflects the harder ResNet18 task rather than the λ setting, and we read it as an efficiency-oriented operating point.

Fig. 1 summarises the CIFAR-100 behaviour. MAML-Select stays close to FedAvg and FedGCS on accuracy and loss while remaining in the low-energy cluster, and panel (b) shows the selector meta-update settling.

Fig. 2(a) shows the efficiency and accuracy trade-off. Relative to FedAvg over shared seeds, MAML-Select lowers cumulative TFLOPs by 12.8%, 21.5% and 1.7% and modelled energy by 16.4%, 24.7% and 4.7% on the three datasets, for accuracy changes of -0.1 , -3.8 and -1.5 percentage points. The main empirical gain is therefore resource efficiency at competitive accuracy.

Participation fairness uses Jain’s index on the per-client selection counts $\{p_i\}_{i=1}^N$, $J = (\sum_i p_i)^2 / (N \sum_i p_i^2)$, which lies in $[1/N, 1]$ and equals 1 for perfectly even selection. Participation coverage $\text{Cov} = \frac{100}{N} |\{i : \sum_t a_{i,t} \geq 1\}|$ is the percentage of clients selected at least once. MAML-Select reaches 100% coverage on every dataset, with a Jain index of 0.61 to 0.78. The two answer different questions. Coverage is guaranteed by the exploration slot (Proposition 2) independently of the learned ranking, whereas the Jain index measures how evenly the remaining $K - \varepsilon$ places are spread. A latency-aware selector uses fast devices more often, and Sec. V-B shows this is what the lower index reflects.

B. Tier-Level Participation

Coverage says how evenly clients are used but not which ones, so we also measured the share of selections per device tier on CIFAR-100 against a pool composition of 0.20/0.50/0.30. FedCS collapses onto the fast tier, placing all of its selections in Tier 3 and reaching only 10% coverage. FedAvg tracks the pool at 0.196/0.495/0.309 and FedGCS at 0.217/0.472/0.311, both with full coverage. MAML-Select




Fig. 2. (a) Efficiency and accuracy trade-off. The x -axis is cumulative training cost relative to FedAvg on a log scale, where $x < 1$ is cheaper than FedAvg, the y -axis is final accuracy, and marker shape encodes the dataset. (b) Measured per-round selection overhead, reported as a mean with 95% confidence intervals over 10 rounds. (c) The analytical operation count $C(N, K) = NP + N \log K + 2KP$ from (11) with $P = 4,673$, normalized to its value at $N = 20$. Panels (b) and (c) vary the pool size N with the selection rate held at 10%.

shifts toward the fast tier at 0.116/0.446/0.437, yet keeps the slow tier represented and retains 100% coverage. The distinction that matters is between preferring fast devices and excluding slow ones, and only the former is compatible with Proposition 2.

C. Coverage Without the Forced Cold Start

The cold start visits every client once in the first 10 of 200 rounds, so the coverage figures could be an artifact of that schedule. Recomputing both metrics over the policy-driven rounds alone, across all 77 MAML-Select runs and every λ , inner-step, width and feature-ablation setting, coverage is 100% in every individual run. Jain’s index falls only slightly, from 0.755 to 0.737, 0.420 to 0.377, and 0.796 to 0.779 on the three datasets. The following result shows why.

Proposition 2 (Coverage from the exploration slot alone). *Suppose every client is available in every round and $\varepsilon \geq 1$. Then, after the cold start, every client is selected at least once in any window of N consecutive rounds, independently of the policy h_ϕ .*

Proof. Staleness obeys $\tau_{i,t+1} = 0$ if $i \in \mathcal{S}_t$ and $\tau_{i,t+1} = \tau_{i,t} + 1$ otherwise, and the slot goes to a client of maximal staleness. Fix j and a window starting at t , and suppose j is never selected in it. For j to miss the slot at round $s \geq t$, some $i_s \neq j$ with $\tau_{i_s,s} \geq \tau_{j,s}$ takes it and resets to 0. Since i_s resets at s , its staleness at $s' > s$ is $s' - s - 1$, so blocking j again would need $s' - s - 1 \geq \tau_{j,t} + (s' - t)$, that is $t - s \geq \tau_{j,t} + 1 \geq 1$, which is impossible for $s \geq t$. Each of the other $N - 1$ clients therefore blocks j at most once, so j is selected within $N - 1$ rounds.

With $N = 100$ and $T = 200$ the condition is met twice over, which is why coverage holds in every run rather than on average. The cold start seeds the feedback buffer but does not

TABLE II

BENCHMARK SUMMARY ACROSS NON-IID DATASETS. ACCURACY, F1, AND COVERAGE ARE IN PERCENT, AND ENERGY IS IN WATT-HOUR (WH). PRECISION AND RECALL ARE OMITTED BECAUSE RECALL COINCIDES WITH ACCURACY UNDER MICRO-AVERAGING. VALUES ARE THE MEAN $\pm$ STANDARD DEVIATION OVER SEEDS.

Method	Acc.	F1	TFLOPs	Energy	Jain	Cov.
<i>Fashion-MNIST</i>						
FedAvg	90.22 $\pm$ 0.58	90.22 $\pm$ 0.51	140 $\pm$ 1	5752 $\pm$ 72	0.95 $\pm$ 0.00	100 $\pm$ 0
FedCS	84.11 $\pm$ 1.82	83.91 $\pm$ 1.98	76 $\pm$ 7	2772 $\pm$ 211	0.10 $\pm$ 0.00	10 $\pm$ 0
Oort	86.70 $\pm$ 0.19	86.42 $\pm$ 0.26	149 $\pm$ 8	5963 $\pm$ 456	0.10 $\pm$ 0.00	10 $\pm$ 0
TiFL	86.73 $\pm$ 0.34	86.51 $\pm$ 0.34	147 $\pm$ 6	6062 $\pm$ 348	0.11 $\pm$ 0.00	12 $\pm$ 0
FedCor	89.35 $\pm$ 0.50	89.22 $\pm$ 0.57	121 $\pm$ 5	4845 $\pm$ 204	0.26 $\pm$ 0.04	49 $\pm$ 8
CriticalFL	90.16 $\pm$ 0.47	90.11 $\pm$ 0.42	304 $\pm$ 143	12552 $\pm$ 5860	0.98 $\pm$ 0.01	100 $\pm$ 0
FedGCS	90.65 $\pm$ 0.29	90.63 $\pm$ 0.26	136 $\pm$ 4	5595 $\pm$ 113	0.95 $\pm$ 0.02	100 $\pm$ 0
MAML-Select	90.11$\pm$0.47	90.04$\pm$0.30	122$\pm$2	4809$\pm$79	0.77$\pm$0.06	100$\pm$0
<i>CIFAR-10</i>						
FedAvg	79.43 $\pm$ 1.84	79.36 $\pm$ 1.78	8341 $\pm$ 52	4798 $\pm$ 24	0.95 $\pm$ 0.00	100 $\pm$ 0
FedCS	50.45 $\pm$ 2.57	49.87 $\pm$ 2.33	4081 $\pm$ 600	2073 $\pm$ 262	0.10 $\pm$ 0.00	10 $\pm$ 0
Oort	60.19 $\pm$ 5.18	59.69 $\pm$ 4.61	8303 $\pm$ 1923	4677 $\pm$ 856	0.10 $\pm$ 0.00	10 $\pm$ 0
TiFL	61.93 $\pm$ 5.04	61.36 $\pm$ 4.60	8200 $\pm$ 2143	4800 $\pm$ 1203	0.11 $\pm$ 0.00	12 $\pm$ 0
FedCor	67.20 $\pm$ 5.33	66.81 $\pm$ 5.37	7748 $\pm$ 695	4452 $\pm$ 435	0.20 $\pm$ 0.02	31 $\pm$ 5
CriticalFL	77.34 $\pm$ 5.21	77.46 $\pm$ 5.21	14239 $\pm$ 4864	8205 $\pm$ 2807	0.98 $\pm$ 0.01	100 $\pm$ 0
FedGCS	77.88 $\pm$ 1.75	77.69 $\pm$ 1.54	7656 $\pm$ 278	4428 $\pm$ 154	0.90 $\pm$ 0.03	100 $\pm$ 0
MAML-Select	75.63$\pm$1.43	75.25$\pm$1.34	6549$\pm$118	3610$\pm$60	0.61$\pm$0.01	100$\pm$0
<i>CIFAR-100</i>						
FedAvg	59.66 $\pm$ 0.29	59.13 $\pm$ 0.40	6331 $\pm$ 6	3626 $\pm$ 18	0.94 $\pm$ 0.00	100 $\pm$ 0
FedCS	27.52 $\pm$ 0.29	26.39 $\pm$ 0.09	5334 $\pm$ 23	2659 $\pm$ 12	0.10 $\pm$ 0.00	10 $\pm$ 0
Oort	29.41 $\pm$ 1.73	28.13 $\pm$ 1.73	5955 $\pm$ 189	3421 $\pm$ 181	0.10 $\pm$ 0.00	10 $\pm$ 0
TiFL	28.52 $\pm$ 5.08	27.22 $\pm$ 5.37	5961 $\pm$ 35	3480 $\pm$ 27	0.11 $\pm$ 0.00	12 $\pm$ 0
FedCor	30.86 $\pm$ 4.70	29.39 $\pm$ 5.00	6473 $\pm$ 64	3663 $\pm$ 8	0.12 $\pm$ 0.00	14 $\pm$ 1
CriticalFL	45.37 $\pm$ 6.13	44.29 $\pm$ 6.75	10235 $\pm$ 1354	5879 $\pm$ 779	0.98 $\pm$ 0.01	100 $\pm$ 0
FedGCS	58.66 $\pm$ 0.11	58.08 $\pm$ 0.20	6140 $\pm$ 24	3526 $\pm$ 4	0.82 $\pm$ 0.02	100 $\pm$ 0
MAML-Select	58.15$\pm$0.51	57.66$\pm$0.33	6224$\pm$20	3457$\pm$6	0.78$\pm$0.03	100$\pm$0

create the result, and Jain’s index depends mildly on it, so we report both figures.

D. Selector Convergence and Objective Drift

We instrument a full 199-round run on each dataset, recording the inner-step change, the realized query objective, and the meta-update magnitude. Lemma 1 is confirmed exactly, since the inner step is non-increasing on the support loss in 198 of 198 adapted rounds on all three datasets, so $\beta = 10^{-2}$ lies inside the admissible range in practice.

The query objective is not monotone on any dataset. On Fashion-MNIST it falls from 3.312 to 0.015 with round-to-round jumps as large as 2.672, on CIFAR-10 it moves within $[0.015, 1.023]$ with jumps to 0.703, and on CIFAR-100 within $[0.055, 3.704]$ with jumps to 3.612. A fixed objective descended with $\eta \leq 1/L_q$ cannot produce increases of this size, so the sequence is direct evidence that q_t changes between rounds, which is the drift term V_T of Corollary 1. In line with Remark 1, the meta-update magnitude in Fig. 1(b) settles at a small non-zero level rather than decaying to zero. Logging the two consecutive objectives at a common iterate measures V_T directly. Over 198 rounds it is -2.89 on Fashion-MNIST, 0.51 on CIFAR-10, and 1.79 on CIFAR-100, so $|V_T|/T$ is at most 0.015. The partial sums flatten instead of growing with the horizon, which places all three runs in the $V_T = o(T)$ regime of Remark 1, and Fashion-MNIST is negative and so falls in the benign case that remark identifies. The drift term is therefore small but not zero, which is the regime the corollary was written for. These are single instrumented runs and we read them as a measurement of the drift term rather than as a distributional claim.




Fig. 3. Sensitivity to the cost trade-off λ on Fashion-MNIST (left) and CIFAR-10 (right). The black curve is accuracy. Computational cost, energy consumption, and fairness are normalized to their $\lambda = 0.1$ values. The trade-off is mild on Fashion-MNIST and sharper on CIFAR-10.

E. Sensitivity and Ablation

These experiments test whether MAML-Select stays stable when its design choices are perturbed. We probe the cost trade-off λ , the inner-loop steps, the selector width, and the six-feature state, each on the dataset where it is most informative.

Fig. 3 reports the λ sweep. A larger λ weights cheaper clients more, so cost and energy fall on both datasets. Accuracy stays within about one point on Fashion-MNIST, whereas on CIFAR-10 the trade-off is sharper, falling from 64.9% at $\lambda = 0.1$ to 56.3% at $\lambda = 5$ while cost drops by about 26% and energy by about 30%. Fairness also falls at large λ . We keep $\lambda = 0.5$, which stays close to the best accuracy, lowers cost, and avoids the sharp fairness loss seen at $\lambda \geq 1$.

Table III depicts the choices of two ablations. On CIFAR-10, a single inner step gives the highest accuracy (63.8%) and the lowest overhead (38.7 versus 45.2 ms at five steps), so we use one step. On Fashion-MNIST, accuracy stays within about 0.6 points as the width grows over $h \in \{16, 32, 64, 128\}$

(401 to 17,537 parameters), so the compact $h = 64$ default is robust.

TABLE III

DESIGN ABLATIONS, GROUPED BY THE CHOICE BEING VARIED. ACCURACY IS THE MEAN $\pm$ STANDARD DEVIATION OVER THREE SEEDS. THE COST COLUMN REPORTS SELECTION OVERHEAD IN MILLISECONDS FOR THE FIRST TWO GROUPS AND CUMULATIVE TFLOPS FOR THE STATE-FEATURE GROUP, AND A BLANK MARKS A QUANTITY NOT RECORDED FOR THAT GROUP. DEFAULTS ARE SHOWN IN BOLD.

Setting	Acc. (%)	Cost	Jain
<i>Inner-loop steps, CIFAR-10, cost in ms</i>			
1	63.8$\pm$1.8	38.7	
2	63.1 $\pm$ 1.8	39.8	
5	62.2 $\pm$ 2.4	45.2	
<i>Selector width h, Fashion-MNIST, cost in ms</i>			
16	90.0 $\pm$ 0.3	10.2	
32	90.6 $\pm$ 0.3	10.4	
64	90.1$\pm$0.5	10.2	
128	90.3 $\pm$ 0.7	11.6	
<i>State features, Fashion-MNIST, cost in TFLOPs</i>			
Full state vector	90.23$\pm$0.55	122	0.78
w/o loss	90.16 $\pm$ 0.66	123	0.79
w/o gradient norm	90.30 $\pm$ 0.24	122	0.78
w/o latency	90.14 $\pm$ 0.14	119	0.70
w/o battery	89.90 $\pm$ 0.44	122	0.79
w/o frequency	90.47 $\pm$ 0.13	115	0.65
w/o staleness	90.37 $\pm$ 0.19	123	0.78

The last group of Table III removes one state feature at a time. Accuracy changes stay within about 0.35 points, so the policy is robust to any single removal. The largest drop comes from the battery feature (-0.33 pp), with latency and loss also contributing. Removing the frequency feature raises accuracy slightly but lowers the Jain index the most, from 0.78 to 0.65, so the participation features chiefly protect fairness. These results support the full six-feature state because it preserves both accuracy stability and participation balance.

VI. CONCLUSION AND FUTURE WORK

We introduced MAML-Select, an online adaptive client-selection framework for heterogeneous FL that learns a lightweight ranking policy over both learning feedback and system-state information, adapting it with recent support-set feedback before each selection decision. Across Fashion-MNIST, CIFAR-10, and CIFAR-100, MAML-Select achieves competitive accuracy while reducing cumulative computational cost and modelled energy consumption compared with higher-cost selection strategies. The λ -sensitivity study shows a mild trade-off on Fashion-MNIST and a sharper one on CIFAR-10, so λ is tuned per task, and the ablations confirm that a single inner step, a compact policy network, and the full six-feature state are sufficient for stable selection. Future work will broaden the evaluation across heterogeneity levels, client-pool sizes, and naturally federated benchmarks, and will study an asynchronous extension that relaxes the synchronous-round assumption together with a differentially private treatment of the client-state vector so that selection does not leak device information.

REFERENCES

- [1] B. McMahan *et al.*, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. AISTATS*. PMLR, 2017, pp. 1273–1282.
- [2] S. P. Karimireddy *et al.*, “SCAFFOLD: Stochastic controlled averaging for federated learning,” in *Proc. ICML*. PMLR, 2020, pp. 5132–5143.
- [3] V. S. Mai, R. J. La, and T. Zhang, “A study of enhancing federated learning on non-IID data with server learning,” *IEEE Transactions on Artificial Intelligence*, vol. 5, no. 11, pp. 5589–5604, 2024.
- [4] R. Tapwal, “Stragglers reimagined: Explainability-driven adaptive federated learning for resource constrained IoMT system,” *IEEE Transactions on Artificial Intelligence*, vol. 7, no. 2, pp. 1002–1011, 2026.
- [5] A. Forootani and R. Iervolino, “Asynchronous federated learning with nonconvex client objective functions and heterogeneous dataset,” *IEEE Transactions on Artificial Intelligence*, vol. 7, no. 5, pp. 2610–2624, 2026.
- [6] K. Borazjani, P. Abdisarabshali, N. Khosravan, and S. Hosseinalipour, “Redefining non-IID data in federated learning for computer vision tasks: Migrating from labels to embeddings for task-specific data distributions,” *IEEE Transactions on Artificial Intelligence*, 2026, early Access.
- [7] T. Nishio and R. Yonetani, “Client selection for federated learning with heterogeneous resources in mobile edge,” in *Proc. IEEE ICC*. IEEE, 2019, pp. 1–7.
- [8] Z. Chai *et al.*, “TiFL: A tier-based federated learning system,” in *Proc. ACM HPDC*, 2020, pp. 125–136.
- [9] F. Lai, X. Zhu, H. V. Madhyastha, and M. Chowdhury, “Oort: Efficient federated learning via guided participant selection,” in *Proc. USENIX OSDI*, 2021, pp. 19–35.
- [10] M. Tang *et al.*, “FedCor: Correlation-based active client selection strategy for heterogeneous federated learning,” in *Proc. IEEE/CVF CVPR*. IEEE, 2022, pp. 10 102–10 111.
- [11] Y. Xu *et al.*, “Federated learning with client selection and gradient compression in heterogeneous edge systems,” *IEEE Transactions on Mobile Computing*, vol. 23, no. 5, pp. 5446–5461, 2024.
- [12] G. Yan, H. Wang, X. Yuan, and J. Li, “CriticalFL: A critical learning periods augmented client selection framework for efficient federated learning,” in *Proc. ACM SIGKDD*, 2023, pp. 5034–5044.
- [13] Z. Ning *et al.*, “FedGCS: A generative framework for efficient client selection in federated learning via gradient-based optimization,” in *Proceedings of the 33rd International Joint Conference on Artificial Intelligence*, 2024, pp. 4760–4768.
- [14] H. Wang, Z. Kaplan, D. Niu, and B. Li, “Optimizing federated learning on non-IID data with reinforcement learning,” in *IEEE INFOCOM*. IEEE, 2020, pp. 1698–1707.
- [15] Q. Pan, H. Cao, Y. Zhu, J. Liu, and B. Li, “Contextual client selection for efficient federated learning over edge devices,” *IEEE Transactions on Mobile Computing*, vol. 23, no. 6, pp. 6538–6548, 2024.
- [16] K. Kazemi, M. H. Badiei, and H. Kebriaei, “Robust peer-to-peer federated learning with deep reinforcement learning based client selection against data poisoning attacks,” *IEEE Transactions on Artificial Intelligence*, vol. 7, no. 1, pp. 262–271, 2026.
- [17] C. Finn, P. Abbeel, and S. Levine, “Model-agnostic meta-learning for fast adaptation of deep networks,” in *Proc. ICML*. PMLR, 2017, pp. 1126–1135.
- [18] A. Fallah, A. Mokhtari, and A. Ozdaglar, “Personalized federated learning with theoretical guarantees: A model-agnostic meta-learning approach,” *Proc. NeurIPS*, vol. 33, pp. 3557–3568, 2020.
- [19] H. A. Tran, C. Ta, and T. X. Tran, “FedImp: Enhancing federated learning convergence with impurity-based weighting,” *IEEE Transactions on Artificial Intelligence*, vol. 7, no. 3, pp. 1652–1665, 2026.
- [20] S. Arisdakessian, O. A. Wahab, A. Mourad, and H. Otrouk, “Towards vox populi in federated learning: A fair and inclusive client selection framework,” *IEEE Transactions on Artificial Intelligence*, vol. 7, pp. 1997–2011, 2026.
- [21] X. Wu, R. He, Z. Sun, and T. Tan, “A light CNN for deep face representation with noisy labels,” *IEEE Transactions on Information Forensics and Security*, vol. 13, no. 11, pp. 2884–2896, 2018.
- [22] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE CVPR*. IEEE, 2016, pp. 770–778.
- [23] T. Li, A. K. Sahu, A. Talwalkar, and V. Smith, “Federated learning: Challenges, methods, and future directions,” *IEEE Signal Processing Magazine*, vol. 37, no. 3, pp. 50–60, 2020.